



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('FastagFraudDetection.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geo
0	1	1/6/2023 11:20	Bus	FTG-001-ABC-121	A-101	Express	Large	350	120	1:
1	2	1/7/2023 14:55	Car	FTG-002-XYZ-451	B-102	Regular	Small	120	100	1:
2	3	1/8/2023 18:25	Motorcycle	NaN	D-104	Regular	Small	0	0	1:
3	4	1/9/2023 2:05	Truck	FTG-044-LMN-322	C-103	Regular	Large	350	120	1:
4	5	1/10/2023 6:35	Van	FTG-505-DEF-652	B-102	Express	Medium	140	100	1:

```
In [4]: df.tail()
```

```
Out[4]:
```

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geo
4995	4996	1/1/2023 22:18	Truck	FTG-445-EDC-765	C-103	Regular	Large	330	330	1:
4996	4997	1/17/2023 13:43	Van	FTG-446-LMK-432	B-102	Express	Medium	125	125	1:
4997	4998	2/5/2023 5:08	Sedan	FTG-447-PLN-109	A-101	Regular	Medium	115	115	1:
4998	4999	2/20/2023 20:34	SUV	FTG-458-VFR-876	B-102	Express	Large	145	145	1:
4999	5000	3/10/2023 0:59	Bus	FTG-459-WSX-543	C-103	Regular	Large	330	125	1:

```
In [5]: df.isnull().sum()
```

Out[5]: Transaction_ID 0
Timestamp 0
Vehicle_Type 0
FastagID 549
TollBoothID 0
Lane_Type 0
Vehicle_Dimensions 0
Transaction_Amount 0
Amount_paid 0
Geographical_Location 0
Vehicle_Speed 0
Vehicle_Plate_Number 0
Fraud_indicator 0
dtype: int64

In [6]: df.describe()

Out[6]:

	Transaction_ID	Transaction_Amount	Amount_paid	Vehicle_Speed
count	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	161.06200	141.261000	67.851200
std	1443.520003	112.44995	106.480996	16.597547
min	1.000000	0.00000	0.000000	10.000000
25%	1250.750000	100.00000	90.000000	54.000000
50%	2500.500000	130.00000	120.000000	67.000000
75%	3750.250000	290.00000	160.000000	82.000000
max	5000.000000	350.00000	350.000000	118.000000

In [7]: df.dtypes

Out[7]: Transaction_ID int64
Timestamp object
Vehicle_Type object
FastagID object
TollBoothID object
Lane_Type object
Vehicle_Dimensions object
Transaction_Amount int64
Amount_paid int64
Geographical_Location object
Vehicle_Speed int64
Vehicle_Plate_Number object
Fraud_indicator object
dtype: object

In [8]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 13 columns):
Column Non-Null Count Dtype

0 Transaction_ID 5000 non-null int64
1 Timestamp 5000 non-null object
2 Vehicle_Type 5000 non-null object
3 FastagID 4451 non-null object
4 TollBoothID 5000 non-null object
5 Lane_Type 5000 non-null object
6 Vehicle_Dimensions 5000 non-null object
7 Transaction_Amount 5000 non-null int64
8 Amount_paid 5000 non-null int64
9 Geographical_Location 5000 non-null object
10 Vehicle_Speed 5000 non-null int64
11 Vehicle_Plate_Number 5000 non-null object
12 Fraud_indicator 5000 non-null object
dtypes: int64(4), object(9)
memory usage: 507.9+ KB

In [9]: df.nunique()

Out[9]: Transaction_ID 5000
Timestamp 4423
Vehicle_Type 7
FastagID 4451
TollBoothID 6
Lane_Type 2
Vehicle_Dimensions 3
Transaction_Amount 20
Amount_paid 23
Geographical_Location 5
Vehicle_Speed 85
Vehicle_Plate_Number 5000
Fraud_indicator 2
dtype: int64

```
In [10]: df.shape
```

```
Out[10]: (5000, 13)
```

```
In [11]: df.duplicated().sum()
```

```
Out[11]: 0
```

```
In [12]: import warnings
warnings.filterwarnings('ignore')
```

```
In [13]: df['FastagID'].unique()
```

```
Out[13]: array(['FTG-001-ABC-121', 'FTG-002-XYZ-451', nan, ..., 'FTG-447-PLN-109',
               'FTG-458-VFR-876', 'FTG-459-WSX-543'], dtype=object)
```

```
In [14]: df['FastagID'].value_counts()
```

```
Out[14]: FastagID
FTG-001-ABC-121    1
FTG-524-CDE-098    1
FTG-531-ZAS-987    1
FTG-530-LP0-210    1
FTG-528-WSX-876    1
..
FTG-414-HIJ-567    1
FTG-647-KLM-890    1
FTG-880-NOP-123    1
FTG-113-QRS-456    1
FTG-459-WSX-543    1
Name: count, Length: 4451, dtype: int64
```


```
In [15]: df['Fraud_indicator'].unique()
```

```
Out[15]: array(['Fraud', 'Not Fraud'], dtype=object)
```

```
In [16]: df['Fraud_indicator'].value_counts()
```

```
Out[16]: Fraud_indicator
Not Fraud    4017
Fraud         983
Name: count, dtype: int64
```

```
In [17]: plt.figure(figsize=(15, 6))
sns.countplot(x='Fraud_indicator', data=df, palette='hls')
plt.show()
```



```
In [18]: df.dropna(inplace=True)
```

```
In [19]: df.head()
```

Out[19]:

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geo	
	0	1	1/6/2023 11:20	Bus	FTG-001-ABC-121	A-101	Express	Large	350	120	1:
	1	2	1/7/2023 14:55	Car	FTG-002-XYZ-451	B-102	Regular	Small	120	100	1:
	3	4	1/9/2023 2:05	Truck	FTG-044-LMN-322	C-103	Regular	Large	350	120	1:
	4	5	1/10/2023 6:35	Van	FTG-505-DEF-652	B-102	Express	Medium	140	100	1:
	5	6	1/11/2023 10:00	Sedan	FTG-066-GHI-987	A-101	Regular	Medium	160	100	1:

In [20]:

df.shape

Out[20]:

(4451, 13)

In [21]:

df.isnull().sum()

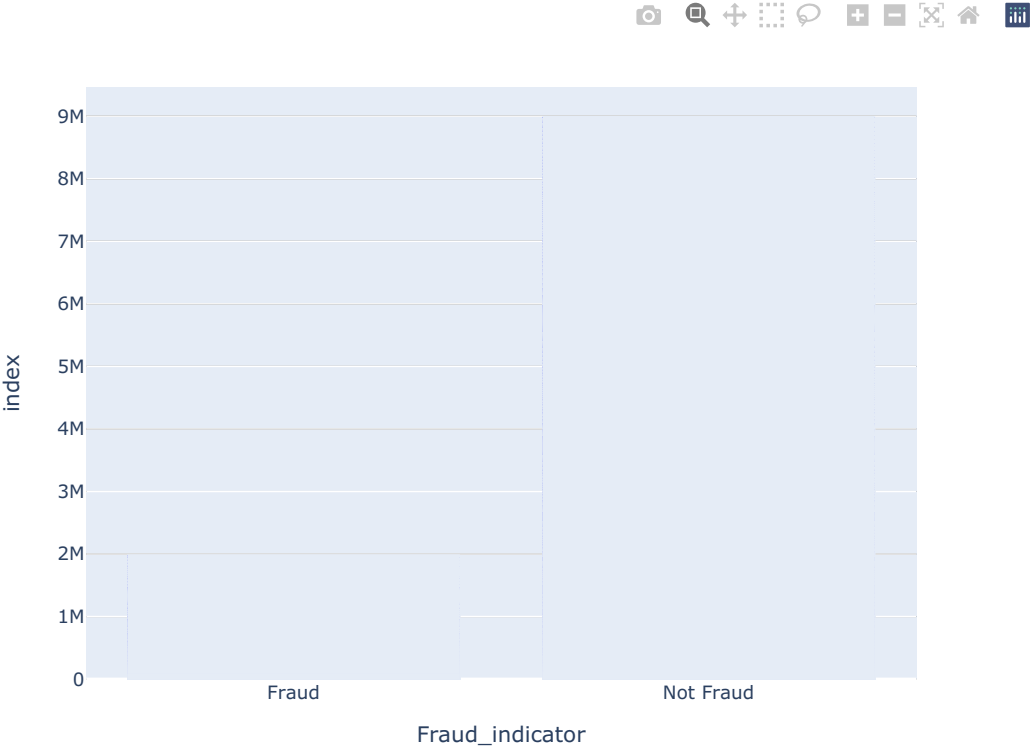
Out[21]:

Transaction_ID0
Timestamp0
Vehicle_Type0
FastagID0
TollBoothID0
Lane_Type0
Vehicle_Dimensions0
Transaction_Amount0
Amount_paid0
Geographical_Location0
Vehicle_Speed0
Vehicle_Plate_Number0
Fraud_indicator0
dtype: int64

In [22]:

import plotly.express as px

fig = px.bar(df, x="Fraud_indicator", y= df.index)
fig.show()



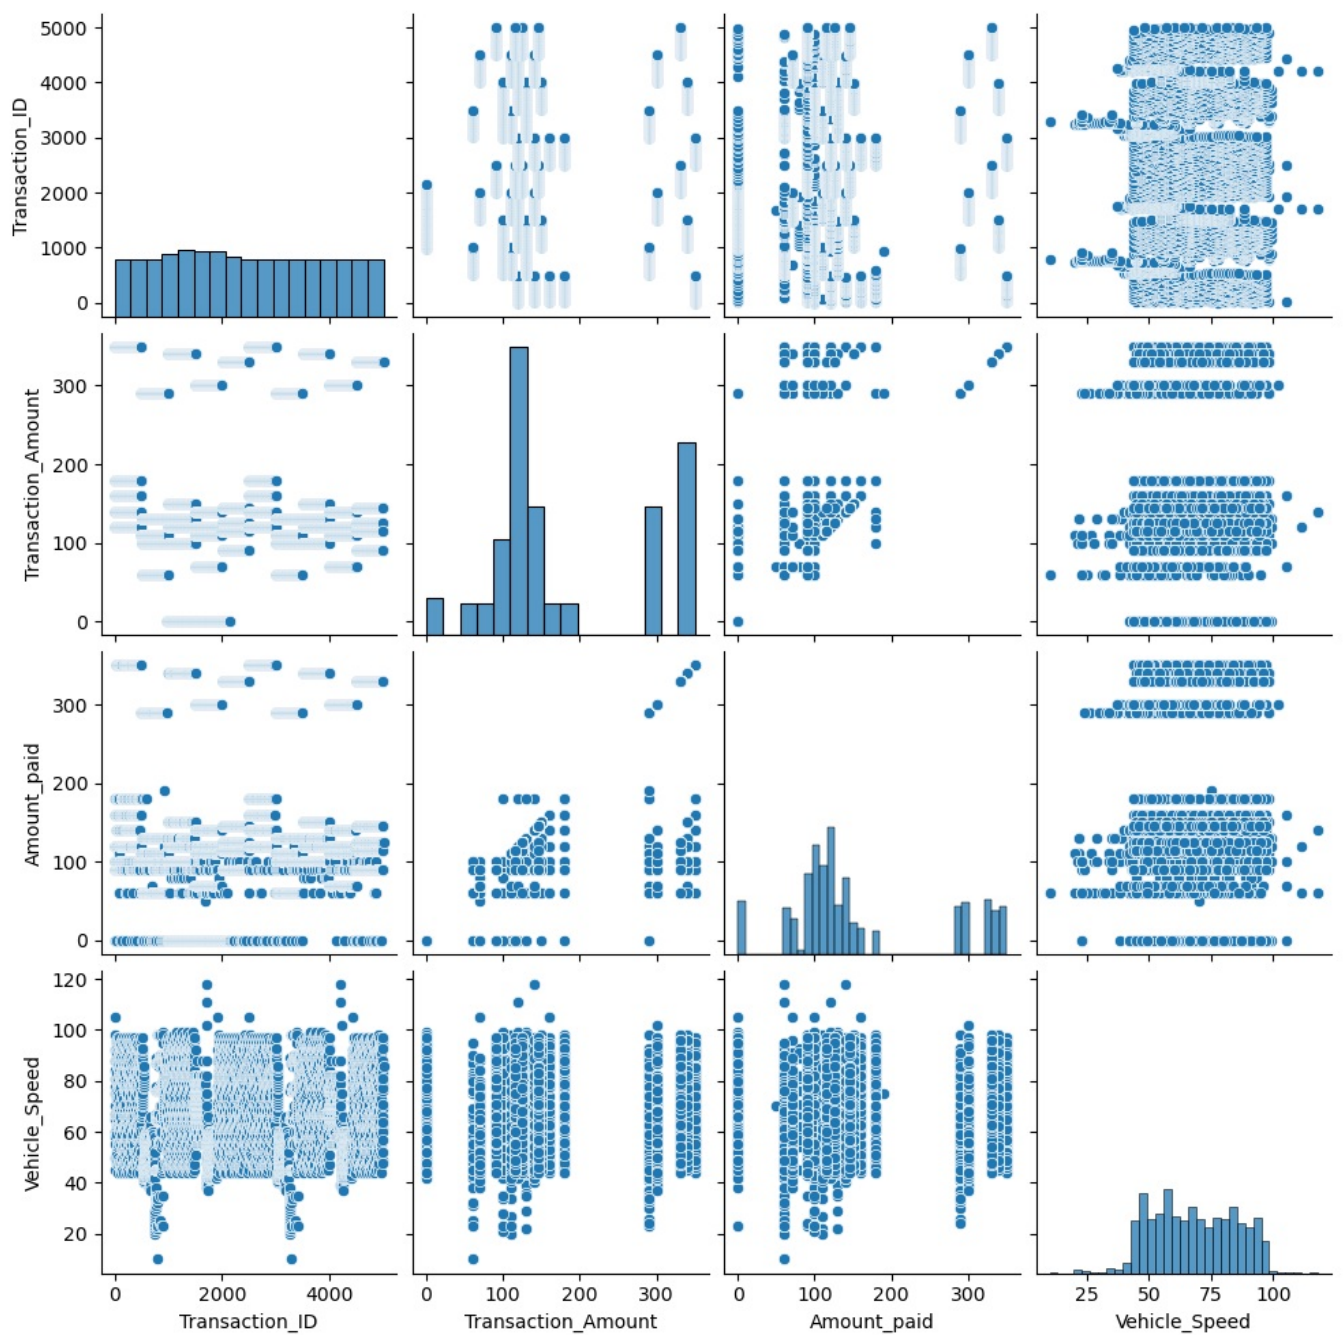
In [23]:

value_counts = df['Fraud_indicator'].value_counts()
fig = px.pie(names=value_counts.index, values=value_counts.values)
fig.update_layout(
 title='Pie Chart of Fraud_indicator',
 title_x=0.5
)

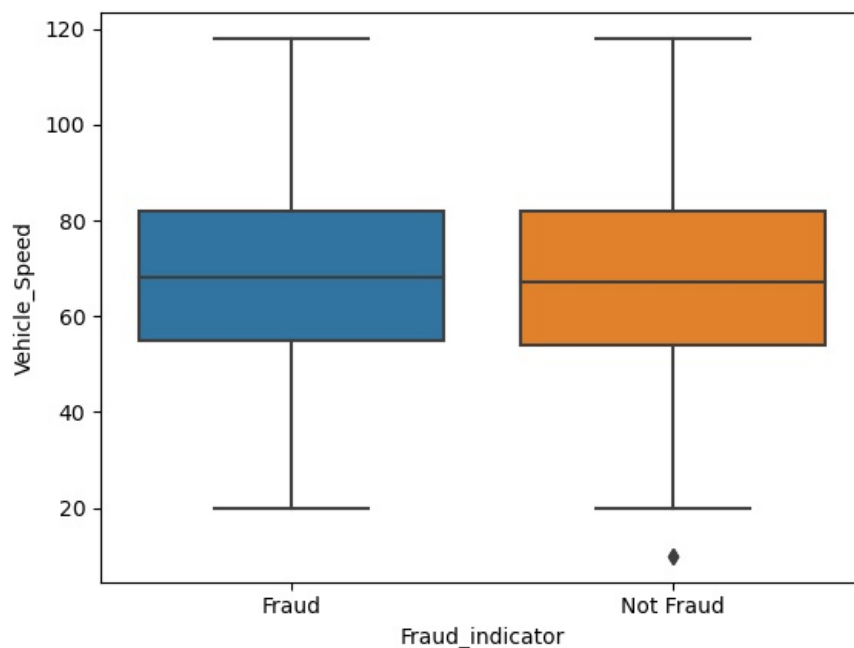
```
fig.show()
```

```
In [24]: sns.pairplot(df)
```

```
Out[24]: <seaborn.axisgrid.PairGrid at 0x1d5b0112a90>
```



```
In [33]: sns.boxplot(x='Fraud_indicator', y='Vehicle_Speed', data=df)
plt.show()
```



```
In [35]: df[['Transaction_Amount', 'Amount_paid', 'Vehicle_Speed']].corr()
```

```
Out[35]:
```

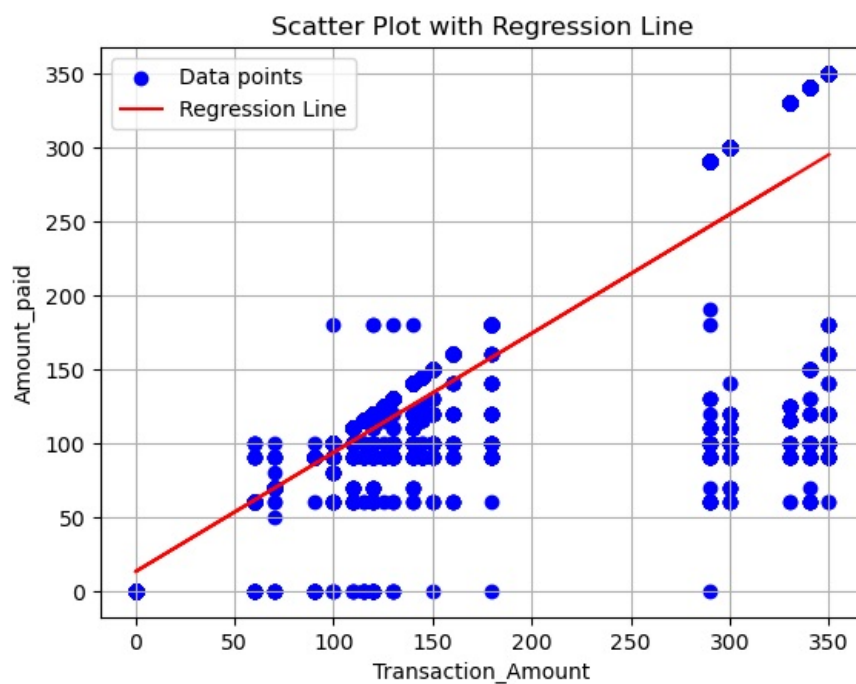
	Transaction_Amount	Amount_paid	Vehicle_Speed
Transaction_Amount	1.000000	0.831275	0.061599
Amount_paid	0.831275	1.000000	0.043446
Vehicle_Speed	0.061599	0.043446	1.000000

```
In [38]: correlation = df['Transaction_Amount'].corr(df['Amount_paid'])
correlation
```


```
Out[38]: 0.8312749747685372
```

```
In [40]: # Calculate the regression line
coefficients = np.polyfit(df['Transaction_Amount'], df['Amount_paid'], 1)
regression_line = np.polyval(coefficients, df['Transaction_Amount'])

# Plot a scatter plot with the regression line
plt.scatter(df['Transaction_Amount'], df['Amount_paid'], color='blue', label='Data points')
plt.plot(df['Transaction_Amount'], regression_line, color='red', label='Regression Line')
plt.title('Scatter Plot with Regression Line')
plt.xlabel('Transaction_Amount')
plt.ylabel('Amount_paid')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [41]: sns.regplot(x='Transaction_Amount', y='Amount_paid', data=df)
plt.show()
```



```
In [43]: df = df.copy()
```

```
df.head(3)
```

```
Out[43]:
```

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geo
0	1	1/6/2023 11:20	Bus	FTG-001-ABC-121	A-101	Express	Large	350	120	1:
1	2	1/7/2023 14:55	Car	FTG-002-XYZ-451	B-102	Regular	Small	120	100	1:
3	4	1/9/2023 2:05	Truck	FTG-044-LMN-322	C-103	Regular	Large	350	120	1:

```
In [44]: X = df[['Transaction_Amount', 'Amount_paid']]  
y = df['Fraud_indicator']
```

```
In [45]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [46]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

```
In [47]: from sklearn.preprocessing import LabelEncoder  
label_encoder = LabelEncoder()  
y_train = label_encoder.fit_transform(y_train)  
y_test = label_encoder.transform(y_test)
```

```
In [48]: df.head()
```

```
Out[48]:
```

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geo
0	1	1/6/2023 11:20	Bus	FTG-001-ABC-121	A-101	Express	Large	350	120	1:
1	2	1/7/2023 14:55	Car	FTG-002-XYZ-451	B-102	Regular	Small	120	100	1:
3	4	1/9/2023 2:05	Truck	FTG-044-LMN-322	C-103	Regular	Large	350	120	1:
4	5	1/10/2023 6:35	Van	FTG-505-DEF-652	B-102	Express	Medium	140	100	1:
5	6	1/11/2023 10:00	Sedan	FTG-066-GHI-987	A-101	Regular	Medium	160	100	1:

```
In [49]: #Neural network model  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(32, activation='relu', input_shape=(X_train.shape[1],)))  
model.add(Dense(16, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

WARNING:tensorflow:From C:\Users\Acer\anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

WARNING:tensorflow:From C:\Users\Acer\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

```
In [50]: model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32)	96
dense_1 (Dense)	(None, 16)	528
dense_2 (Dense)	(None, 1)	17

=====
Total params: 641 (2.50 KB)
Trainable params: 641 (2.50 KB)
Non-trainable params: 0 (0.00 Byte)

```
In [51]: model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

WARNING:tensorflow:From C:\Users\Acer\anaconda3\Lib\site-packages\keras\src\optimizers_init_.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

```
In [52]: model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

Epoch 1/10

WARNING:tensorflow:From C:\Users\Acer\anaconda3\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.nn.RaggedTensorValue is deprecated. Please use tf.nn.compat.v1.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\Acer\anaconda3\Lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

89/89 [=====] - 3s 11ms/step - loss: 0.5047 - accuracy: 0.7788 - val_loss: 0.4317 - val_accuracy: 0.7739

Epoch 2/10

89/89 [=====] - 0s 5ms/step - loss: 0.3707 - accuracy: 0.8294 - val_loss: 0.3243 - val_accuracy: 0.8638

Epoch 3/10

89/89 [=====] - 0s 5ms/step - loss: 0.2643 - accuracy: 0.8912 - val_loss: 0.2259 - val_accuracy: 0.9242

Epoch 4/10

89/89 [=====] - 0s 5ms/step - loss: 0.1900 - accuracy: 0.9305 - val_loss: 0.1768 - val_accuracy: 0.9312

Epoch 5/10

89/89 [=====] - 0s 5ms/step - loss: 0.1507 - accuracy: 0.9501 - val_loss: 0.1464 - val_accuracy: 0.9607

Epoch 6/10

89/89 [=====] - 1s 6ms/step - loss: 0.1239 - accuracy: 0.9579 - val_loss: 0.1288 - val_accuracy: 0.9719

Epoch 7/10

89/89 [=====] - 1s 8ms/step - loss: 0.1063 - accuracy: 0.9733 - val_loss: 0.1185 - val_accuracy: 0.9817

Epoch 8/10

89/89 [=====] - 1s 6ms/step - loss: 0.0952 - accuracy: 0.9810 - val_loss: 0.1096 - val_accuracy: 0.9817

Epoch 9/10

89/89 [=====] - 0s 5ms/step - loss: 0.0862 - accuracy: 0.9814 - val_loss: 0.1058 - val_accuracy: 0.9817

Epoch 10/10

89/89 [=====] - 1s 6ms/step - loss: 0.0800 - accuracy: 0.9821 - val_loss: 0.1014 - val_accuracy: 0.9817

```
Out[52]: <keras.src.callbacks.History at 0x1d5bd481b10>
```

```
In [53]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
y_pred_prob = model.predict(X_test)
```

```
# Convert probabilities to binary predictions
y_pred = np.round(y_pred_prob)
```

```
# Print accuracy metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
# Print accuracy metrics
print("Accuracy: {:.2f}%".format(accuracy * 100))
print("Precision: {:.2f}".format(precision))
print("Recall: {:.2f}".format(recall))
print("F1 Score: {:.2f}".format(f1))
```

28/28 [=====] - 0s 6ms/step

Accuracy: 98.43%

Precision: 0.98

Recall: 1.00

F1 Score: 0.99

```
In [54]: print(y_pred_prob[:5])
y_pred[:5]
```

```
[[0.9332529 ]
 [0.92760545]
 [0.94306105]
 [0.9372957 ]
 [0.16856828]]
Out[54]: array([[1.],
        [1.],
        [1.],
        [1.],
        [0.]], dtype=float32)
```


```
In [55]: import matplotlib.pyplot as plt

metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
values = [accuracy, precision, recall, f1]

# Define custom colors for bars and text
bar_colors = ['skyblue', 'lightgreen', 'lightcoral', 'lightsalmon']
plt.bar(metrics, values, color=bar_colors)

# Adding values on top of each bar
for i, v in enumerate(values):
    plt.text(i, v + 0.05, str(round(v, 2)), ha='center', va='bottom')

plt.ylabel('Score')
plt.title('Model Metrics')
plt.ylim(0, 1.1)
plt.show()
```



```
In [ ]:
```